

schema.py

business world +
unrelated records



records.py

invented names,
 ≥ 3 edits apart



render.py

3-4 templates
per record type



**questions.py
+ lookalike.py**

12 questions per doc
+ their look-alikes



filler.py

unrelated or sibling,
same record count



assemble.py

exact length ($\pm 0.1\%$),
look-alikes at 35-65%



absence_check.py

proof for every
question



pass

build_dataset.py

565 documents,
6,780 questions

fail $\rightarrow$ rebuild with a new seed

